

In [1]:

```

import os
import time
from PIL import Image
import pandas as pd

import pyspark
import os as os
from pyspark.sql import SparkSession

from sklearn.manifold import TSNE
import time
import seaborn as sns

#customised functions
def set_up_working_directory(str_dir_path):
    os.chdir(str_dir_path)
    print("Working direcotry changed to " + os.getcwd())

def data_basic_details(image_folder_path):

    df = pd.DataFrame(columns = ['FileName', 'width', 'height'])
    print(df)
    start_time = time.time()
    #iterate through every file
    for r, d, f in os.walk(image_folder_path):
        for file in f:
            if file.endswith(".jpg"):
                width, height = Image.open(os.path.join(r, file)).size
                df = df.append({'FileName' : file , 'width' :width , 'height' :
#total time taken
print('Time Taken:', time.strftime("%H:%M:%S",time.gmtime(time.time() - start

def spark_spin_up_session(session_name):
    spark = SparkSession.builder.appName(session_name).getOrCreate()
    return spark

def spark_read_csv(sparkobj, file_path):
    spark = sparkobj
    readobj= spark.read.option('header','true').csv(file_path)
    return readobj

def spark_read_images_from_path(sparkobj, file_path):
    spark = sparkobj
    image_df = spark.read.format("image").load(file_path, inferSchema=True)
    return image_df

def spark_return_image_attribute(image_obj):
    image_obj.select("image.origin", "image.width", "image.height", "image.nChannel

```

```
In [2]: spark= spark_spin_up_session("session1")
spark
```

**Out[2]: SparkSession - in-memory
SparkContext**

[Spark UI \(http://Narayan-PC:4040\)](http://Narayan-PC:4040)

Version

v3.2.0

Master

local[*]

AppName

session1

```
In [3]: image_df= spark_read_images_from_path(spark, "e:/dev/Kaggle/fashion/sample/1164.jpg")
spark_return_image_attribute(image_df)
```

origin	width	height	nChannels	mode
file:///e:/dev/Kaggle/fashion/sample/1164.jpg	60	80	3	16

```
In [4]: image_df = spark_read_images_from_path(spark, "e:/dev/Kaggle/fashion/sample/")
spark_return_image_attribute(image_df)
```

origin	width	height	nChannels	mode
file:///e:/dev/Kaggle/fashion/sample/1611.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1619.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1755.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1752.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1533.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1689.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1982.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1570.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1531.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1617.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1981.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1613.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1164.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1603.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1625.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1798.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1536.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1528.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1539.jpg	60	80	3	16
file:///e:/dev/Kaggle/fashion/sample/1636.jpg	60	80	3	16

only showing top 20 rows

In [5]: `image_df.printSchema()`

```
root
|-- image: struct (nullable = true)
|    |-- origin: string (nullable = true)
|    |-- height: integer (nullable = true)
|    |-- width: integer (nullable = true)
|    |-- nChannels: integer (nullable = true)
|    |-- mode: integer (nullable = true)
|    |-- data: binary (nullable = true)
```

In [6]: `display(image_df)`

```
DataFrame[image: struct<origin:string,height:int,width:int,nChannels:int,mode:integer,data:binary>]
```

In [7]: `from mpl_toolkits.mplot3d import Axes3D`
`from sklearn.preprocessing import StandardScaler`
`import matplotlib.pyplot as plt # plotting`
`import matplotlib.image as mpimg`
`import numpy as np # linear algebra`
`import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)`
`import os # accessing directory structure`

In [8]: `DATASET_PATH = "e:/dev/Kaggle/fashion/"`
`print(os.listdir(DATASET_PATH))`

```
['df_embs.txt', 'embed.txt', 'images', 'myntdataset', 'my_model.h5', 'project', 'sample', 'styles.csv']
```

```
In [9]: df = pd.read_csv(DATASET_PATH + "styles.csv", nrows=5000, error_bad_lines=False)
df['image'] = df.apply(lambda row: str(row['id']) + ".jpg", axis=1)
df = df.reset_index(drop=True)
df.head(10)
```

Out[9]:

	id	gender	masterCategory	subCategory	articleType	baseColour	season	year	usage	
0	15970	Men	Apparel	Topwear	Shirts	Navy Blue	Fall	2011	Casual	
1	39386	Men	Apparel	Bottomwear	Jeans	Blue	Summer	2012	Casual	
2	59263	Women	Accessories	Watches	Watches	Silver	Winter	2016	Casual	
3	21379	Men	Apparel	Bottomwear	Track Pants	Black	Fall	2011	Casual	I
4	53759	Men	Apparel	Topwear	Tshirts	Grey	Summer	2012	Casual	
5	1855	Men	Apparel	Topwear	Tshirts	Grey	Summer	2011	Casual	
6	30805	Men	Apparel	Topwear	Shirts	Green	Summer	2012	Ethnic	
7	26960	Women	Apparel	Topwear	Shirts	Purple	Summer	2012	Casual	
8	29114	Men	Accessories	Socks	Socks	Navy Blue	Summer	2012	Casual	
9	30039	Men	Accessories	Watches	Watches	Black	Winter	2016	Casual	



```
In [10]: import cv2
def plot_figures(figures, nrows = 1, ncols=1,figsize=(8, 8)):
    """Plot a dictionary of figures.

    Parameters
    -----
    figures : <title, figure> dictionary
    ncols : number of columns of subplots wanted in the display
    nrows : number of rows of subplots wanted in the figure
    """

    fig, axeslist = plt.subplots(ncols=ncols, nrows=nrows,figsize=figsize)
    for ind,title in enumerate(figures):
        axeslist.ravel()[ind].imshow(cv2.cvtColor(figures[title], cv2.COLOR_BGR2RGB))
        axeslist.ravel()[ind].set_title(title)
        axeslist.ravel()[ind].set_axis_off()
    plt.tight_layout() # optional

def img_path(img):
    return DATASET_PATH+"images/"+img

def load_image(img):
    return cv2.imread(img_path(img))
```

```
In [33]: import matplotlib.pyplot as plt
import numpy as np

# generation of a dictionary of (title, images)
figures = {'image '+str(i): load_image(row.image) for i, row in df.sample(6).iter
# plot of the images in a figure, with 2 rows and 3 columns
plot_figures(figures, 2, 3)
```

image 2746



image 884



image 2495



image 1829



image 3923




image 4319



```
In [46]: plt.figure(figsize=(14,8))  
df.articleType.value_counts().nlargest(15).plot(kind='bar', color="blue")
```

Out[46]: <matplotlib.axes._subplots.AxesSubplot at 0x596d1d60>



```
In [13]: import tensorflow as tf  
from tensorflow.keras import Model  
from tensorflow.keras.applications.resnet50 import ResNet50  
from tensorflow.keras.preprocessing import image  
from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions  
from tensorflow.keras.layers import GlobalMaxPooling2D  
tf.__version__
```

Out[13]: '2.4.0'

```
In [14]: # Input Shape
img_width, img_height, _ = load_image(df.iloc[0].image).shape

# Pre-Trained Model
base_model = ResNet50(weights='imagenet',
                      include_top=False,
                      input_shape = (img_width, img_height, 3))
base_model.trainable = False

# Add Layer Embedding
model = tf.keras.Sequential([
    base_model,
    GlobalMaxPooling2D()
])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
resnet50 (Functional)	(None, 3, 2, 2048)	23587712
=====		
global_max_pooling2d (Global	(None, 2048)	0
=====		
Total params: 23,587,712		
Trainable params: 0		
Non-trainable params: 23,587,712		
=====		

```
In [15]: def get_embedding(model, img_name):
# Reshape
img = image.load_img(img_path(img_name), target_size=(img_width, img_height))
# img to Array
x = image.img_to_array(img)
# Expand Dim (1, w, h)
x = np.expand_dims(x, axis=0)
# Pre process Input
x = preprocess_input(x)
return model.predict(x).reshape(-1)
```

```
In [ ]: def get_embedding(DATASET_PATH, model, img_name, img_width, img_height):
# Reshape
img = image.load_img(img_path(DATASET_PATH, img_name), target_size=(img_width,
# img to Array
x = image.img_to_array(img)
# Expand Dim (1, w, h)
x = np.expand_dims(x, axis=0)
# Pre process Input
x = preprocess_input(x)
return model.predict(x).reshape(-1)
```



```
In [16]: emb = get_embedding(DATASET_PATH, model, df.iloc[1].image)
emb.shape
```

Out[16]: (2048,)

```
In [29]: plt.imshow(cv2.cvtColor(load_image(df.iloc[1].image), cv2.COLOR_BGR2RGB))
print(emb)
```

[1.8793364 1.6008836 0.09203261 ... 3.2688963 2.2717304 5.4333878]



```
In [18]: %%time
import swifter

# Parallel apply
map_embeddings = df['image'].swifter.apply(lambda img: get_embedding(model, img))
df_embs       = map_embeddings.apply(pd.Series)

print(df_embs.shape)
df_embs.head()
```

Pandas Apply: 100%


5000/5000 [07:16<00:00, 11.45it/s]

(5000, 2048)
Wall time: 7min 56s

Out[18]:

	0	1	2	3	4	5	6	7	8	9
0	0.000000	3.539255	0.000000	1.094599	0.000000	0.000000	4.458534	2.446015	2.678130	0.0
1	1.879336	1.600884	0.092033	4.433075	0.000000	0.000000	3.030769	8.530592	5.498659	0.0
2	0.000000	0.311199	0.000000	3.808681	0.437031	7.112498	0.000000	0.000000	0.000000	0.0
3	0.588018	9.894616	0.000000	6.295309	1.783727	1.913123	0.000000	13.309944	7.805779	0.0
4	0.000000	1.882976	0.000000	5.122097	0.000000	0.000000	2.950291	6.626864	1.606633	0.0

5 rows × 2048 columns



```
In [19]: # Load distance metrics
from sklearn.metrics.pairwise import pairwise_distances

#find distance metrics
cosine_sim = 1 - pairwise_distances(df_embs, metric='cosine')
cosine_sim[:4, :4]
```

```
Out[19]: array([[0.99999934, 0.5813052 , 0.23863798, 0.49294078],
               [0.5813052 , 0.99999905, 0.23951322, 0.72273475],
               [0.23863798, 0.23951322, 0.9999998 , 0.22011638],
               [0.49294078, 0.72273475, 0.22011638, 1.          ]], dtype=float32)
```

```
In [52]: indices = pd.Series(range(len(df)), index=df.index)
indices

# Function that get movie recommendations based on the cosine similarity score of
def get_recommender(idx, df, max_rec = 5):
    sim_idx = indices[idx]
    sim_scores = list(enumerate(cosine_sim[sim_idx]))
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)
    sim_scores = sim_scores[1:max_rec+1]
    idx_rec = [i[0] for i in sim_scores]
    idx_sim = [i[1] for i in sim_scores]

    return indices.iloc[idx_rec].index, idx_sim

get_recommender(2993, df, max_rec = 3)
```


```
Out[52]: (Int64Index([259, 4305, 0], dtype='int64'), [0.9040973, 0.8925514, 0.89175576])
```

```
In [53]: # Idx Item to Recommender
idx_ref = 4100

# Recommendations
idx_rec, idx_sim = get_recommender(idx_ref, df, max_rec = 3)

# Plot
#=====
plt.imshow(cv2.cvtColor(load_image(df.iloc[idx_ref].image), cv2.COLOR_BGR2RGB))

# generation of a dictionary of (title, images)
figures = {'im'+str(i): load_image(row.image) for i, row in df.loc[idx_rec].iterrows()}
# plot of the images in a figure, with 2 rows and 3 columns
plot_figures(figures, 1, 3)
```



im4951



im3311



im901



```
In [54]: idx_ref = 987

# Recommendations
idx_rec, idx_sim = get_recommender(idx_ref, df, max_rec = 3)

# Plot
#=====
plt.imshow(cv2.cvtColor(load_image(df.iloc[idx_ref].image), cv2.COLOR_BGR2RGB))

# generation of a dictionary of (title, images)
figures = {'im'+str(i): load_image(row.image) for i, row in df.loc[idx_rec].iterrows()}
# plot of the images in a figure, with 2 rows and 3 columns
plot_figures(figures, 1, 3)
```



im3849

im1779

im1273




```
In [25]: time_start = time.time()
         tsne = TSNE(n_components=2, verbose=0, perplexity=40, n_iter=300)
         tsne_results = tsne.fit_transform(df_embs)
         print('t-SNE done! Time elapsed: {} seconds'.format(time.time()-time_start))
```

t-SNE done! Time elapsed: 119.30700016021729 seconds

```
In [26]: df['tsne-2d-one'] = tsne_results[:,0]
         df['tsne-2d-two'] = tsne_results[:,1]
```

```
In [27]: plt.figure(figsize=(16,10))
         sns.scatterplot(x="tsne-2d-one", y="tsne-2d-two",
                        hue="masterCategory",
                        data=df,
                        legend="full",
                        alpha=0.8)
```

Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x484babe0>



```
In [28]: plt.figure(figsize=(20,16))
sns.scatterplot(x="tsne-2d-one", y="tsne-2d-two",
                hue="subCategory",
                data=df,
                legend="full",
                alpha=0.8)
```

Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x5d1972e0>

